

Model	Method	Cancer-Myth		CREPE	
		TPR ↑	FPR ↓	TPR ↑	FPR ↓
Qwen2.5 7B	Direct + scope protection	—	—	—	—
GPT-6	Hidden probe	75.0 ^H	47.7 ^H	—	—
Luna	Direct + scope protection	—	—	—	—
Gemma 4 12B	Direct + scope protection	—	—	—	—
	Hidden probe	—	—	—	—
Qwen3.8 27B OFF	Direct + scope protection	—	—	—	—
	Hidden probe	—	—	—	—
Qwen3.8 27B ON	Direct + scope protection	—	—	—	—
	Hidden probe	—	—	—	—
Model-independent	TF-IDF text	86.1 ^H	51.0 ^H	—	—

Table S3. Additional detection baselines

Stored results; — means unavailable or not comparable. Values are percentages, one run.

TPR ↑: FPQ detection; FPR ↓: false positives on normal questions.

H: historical full-corpus evaluation (583 FPQ / 149 NFP); valid denominators in ledger.

Default crossfit decisions for probes / TF-IDF, not matched-FPR thresholds.

Model versions, counts, judge/protocol differences and source hashes: result_sources.md.